

Comparative Study of Transfer Learning Architectures for 120-Class Handwritten Javanese Script Recognition Under Severe Class Imbalance

1st Fathan Fardian Sanum

School of Computing
Telkom University
Bandung, Indonesia

fathansanum@student.telkomuniversity.ac.id

2nd Syahdan Rizqi Ruhendy

School of Computing
Telkom University
Bandung, Indonesia
syahdanruhendy@student.
telkomuniversity.ac.id

3rd Ardia Alif Ramadhan

School of Computing
Telkom University
Bandung, Indonesia
ardiarahaman@student.telkomuniversity.ac.id

Abstract—Javanese script (*Hanacaraka*) is a traditional writing system facing the threat of functional extinction due to a declining number of generations able to read and write it, making digitization through automated character recognition an urgent preservation strategy. Two challenges that have not been addressed simultaneously in the literature are full coverage of 120 syllabic classes and extreme class imbalance, with an imbalance ratio of $33.10\times$ (largest class 695 images versus smallest class 21 images) between the Vowel A classes as the majority and the five other vowel groups. This study proposes a four-step structure-preserving data preparation pipeline that explicitly avoids geometric augmentation to maintain the visual integrity of the script, combined with a weighted loss function to compensate for class imbalance. Four transfer learning architectures — EfficientNet-B0, ResNet-18, VGG-16, and MobileNetV3-Large — are evaluated comparatively using an identical pipeline, with metrics reported separately for majority and minority classes. EfficientNet-B0 achieves 97.08% accuracy with a macro F1-score of 0.9110, comparable to studies using artificially balanced datasets, while ResNet-18 exhibits an F1 gap between majority and minority classes of only 0.0004. These results demonstrate that extreme class imbalance can be overcome with a targeted pipeline, and confirm that per-class-group analysis is necessary to reveal model generalization capabilities hidden by aggregate metrics.

Index Terms—*Abugida*, augmented character recognition, class-imbalanced dataset, EfficientNet-B0, handwritten character recognition, Javanese script, MobileNetV3-Large, ResNet-18, structure-preserving augmentation, transfer learning, VGG-16, weighted loss

I. INTRODUCTION

Javanese script (*Hanacaraka*) is an *abugida* writing system combining 20 base consonants with six vowel markers into 120 syllabic classes [1]. The dominance of Latin script in modern education has caused a cross-generational decline in proficiency [2], [3], making digitization via *Optical Character Recognition* (OCR) an urgent preservation intervention [4]. The primary challenge is that the only public dataset covering

all 120 classes, *Indonesian Local Script Characters* [3], exhibits extreme imbalance: Vowel A averages 500 images per class while 100 non-A classes have only around 21 images each, yielding an imbalance ratio of $33.10\times$.

Transfer learning based on CNNs has proven effective for handwritten regional script recognition with limited data [5], [6], significantly outperforming *shallow learning* methods [7]. Combined with a weighted loss function, this approach provides a suitable framework for simultaneously addressing the visual complexity of Javanese script and its class imbalance.

Prior studies leave three gaps not yet addressed simultaneously: Susanto et al. [8] achieved 99.65% accuracy on 120 classes but on an artificially balanced dataset, while Wisesty et al. [3] found that geometric augmentation actually degrades Javanese script recognition performance. These three gaps are: 120-class recognition under natural class imbalance, structurally safe augmentation strategies, and comparative cross-architecture evaluation.

This study closes all three gaps by proposing a four-step structure-preserving data preparation pipeline and evaluating four architectures (EfficientNet-B0, ResNet-18, VGG-16, MobileNetV3-Large) comparatively. EfficientNet-B0 achieves 97.08% accuracy, while ResNet-18 exhibits an F1 gap across vowel groups of only 0.0004, a finding revealed only through per-vowel-group analysis.

II. RELATED WORK

The following studies are the most closely related to this work because they directly target the full 120 syllabic class coverage of Javanese script, serving as the primary benchmarks to surpass. Susanto et al. [9] developed a custom CNN with four convolutional layers and two *fully connected* layers for 120-class classification, achieving 97.29% accuracy on a balanced dataset of 14,400 images from 30 writers. The same group [8] subsequently proposed a 12-layer CNN with

selective augmentation, where rotation was restricted to $\pm 3^\circ$ to avoid altering the visual meaning of the script, achieving 99.65% accuracy. Despite the high accuracy figures, all experiments were conducted on datasets artificially balanced with an equal number of samples per class, so model performance under naturally imbalanced data distributions has never been explored.

Beyond studies specifically targeting 120 classes, several studies covering the broader Nusantara script landscape provide valuable insights into architecture selection, augmentation behavior, and model capability limits on 20-class Javanese script. Wisesty et al. [3] proposed MNetNCR based on MobileNetV3 [10] for five Nusantara scripts and achieved $F1 = 0.979$ for Javanese script without any augmentation. More importantly, applying augmentation to this model actually degraded performance to 0.9579 because geometric transformations destroyed the visual characteristics distinguishing between classes — an empirical warning that became the primary foundation for the augmentation strategy in this study. Nindya et al. [6] applied VGG-16 [11] to multi-script classification of Balinese, Javanese, and Sundanese scripts, achieving 93.19% accuracy for the Javanese subset, and introduced the *Character Likeness Rate* (CLR) metric identifying the “ha”/“la” pair (CLR 25.3%) and “sa”/“da” (CLR 15.8%) as the most confusion-prone characters due to visual similarity. Sulistiyo et al. [7] compared various *shallow learning* methods on the same dataset and found that LightGBM achieved only 77.15% F1-score, far below CNN-based approaches, reinforcing the superiority of deep learning for this task. Ifriza et al. [5] demonstrated that ResNet-18 [12] fine-tuned with a *transfer learning* strategy can achieve 95.53% accuracy from only 1,562 images of 20-class Javanese script, making it the sole ResNet-based *transfer learning* reference for Javanese script prior to this study. All studies in this section, while informative, are limited to balanced datasets and have never tested model robustness under extreme class imbalance.

Class imbalance is a critical aspect overlooked by almost all Javanese script studies, yet it is the natural characteristic of the publicly available dataset. Faizin et al. [2] is the only study to explicitly address this issue, developing SkelBAGAN (*Skeleton-based Balancing GAN*) for a Javanese script dataset with imbalance ratios up to $81.78\times$. SkelBAGAN integrates a character skeleton-based *autoencoder* and *adversarial loss* to generate synthetic images preserving script structure integrity, and demonstrably improves recognition performance on minority classes. Its limitation lies in the GAN architecture complexity requiring significant computational resources, and its scope remaining limited to 20 base classes without testing on the full 120-class syllabic scenario. This study fills that gap with a lighter approach — structure-preserving augmentation and a weighted loss function — without requiring additional generative network training.

To understand the positioning of this study in the broader landscape, a review of OCR studies across Southeast Asian abugida scripts provides context on how far research in this area has advanced. Wicaksono et al. [1] conducted a

TABLE I
CLASS DISTRIBUTION OF DATASET BY VOWEL GROUP

Vowel Group	Classes	Images/Class	Total
Vowel A	20	426–695	9,995
Vowel E	20	21–22	440
Vowel I	20	21–22	440
Vowel O	20	21–22	440
Vowel U	20	21–22	440
Vowel E-taling	20	21–22	440
Total	120	21–695	12,191

IR = $33.10\times$, median = 22.

systematic literature review on six Southeast Asian abugida scripts and found no studies in the past five years addressing word segmentation for either Javanese or Sundanese script, confirming that isolated character recognition remains the frontier of Javanese OCR research. Widiarti and Adji [4] established a baseline evaluation for printed (not handwritten) Javanese OCR, noting that handwriting style variation adds a far more complex layer of challenges compared to printed script. Together, these two studies underscore that the domain of handwritten Javanese character recognition — particularly at 120-class coverage with imbalanced data — remains very open and serves as the direct motivation for this study.

III. DATASET AND EXPERIMENTAL SETUP

This study uses the *Indonesian Local Script Characters* dataset, curated by Aditya Firman Ihsan et al. at Telkom University and available on Mendeley Data, which was also used by prior studies [3], [7]. For the Javanese script subset, the dataset consists of two parts with a total of 13,257 images: the `all_class/` subdirectory containing 12,191 images across 120 syllabic classes (6 vowels $\times$ 20 consonants), and the `variations/` subdirectory containing 1,066 handwriting style variation images for 8 consonants (ba, ga, ha, ka, la, nya, ra, ta) in the Vowel A group. Images predominantly use JPEG format (91%) with RGB color channels (98.3%) and an original average resolution of 362×347 pixels, subsequently standardized to 224×224 pixels as model input.

The class distribution in this dataset is severely imbalanced by nature, as summarized in Table I. The twenty Vowel A classes have 426–695 images per class ($\mu = 499.8$), while 100 non-A classes have only 21–22 images per class, yielding an imbalance ratio of $33.10\times$ (largest/smallest class: 695/21) with a per-class median of only 22, far below the mean of 101.59. This imbalance is the primary challenge that all pipeline stages are designed to address.

The dataset was split per-class using stratified sampling at a 70:15:15 ratio into training (8,483 original images), validation (1,789 images), and test (1,919 images) sets. Per-class stratification ensures all 120 classes are represented in every subset, including minority classes with only 21 images. After augmentation on the training set, the total training samples

Pipeline Persiapan Data Empat Langkah




Fig. 1. Four-step data preparation pipeline applied sequentially before model training.

increased to 38,510 images. All images were converted to 3-channel grayscale and normalized using ImageNet statistics ($\mu = [0.485; 0.456; 0.406]$, $\sigma = [0.229; 0.224; 0.225]$).

IV. PROPOSED METHOD

A. Pipeline Overview

The primary challenge in this study lies not in the model architecture, but in the data conditions — specifically the extreme class imbalance between Vowel A and non-A classes that causes conventional training strategies to produce models systematically biased toward majority classes. To address this, the study proposes a four-step data preparation pipeline executed sequentially before the training process begins. Each step is designed to handle one specific aspect of the data distribution problem: maintaining minority class representation in the split, adding data diversity without new collection, filling minority classes with structurally safe augmentation, and applying proportional penalty weights during training. Fig. 1 illustrates the overall flow of this pipeline.

B. Step 1: Per-Class Stratified Split

The dataset is split into training (70%), validation (15%), and test (15%) sets using per-class stratified sampling. This ensures all 120 classes are represented in every subset, even for minority classes with only 21 images, yielding ≈ 15 training images, 3 validation images, and 3 test images per minority class before augmentation. This split yields 1,789 validation images and 1,919 test images.

C. Step 2: Variation Data Integration

A total of 1,066 images from `variations/` are added to the training set for the 8 corresponding Vowel A consonant classes (ba, ga, ha, ka, la, nya, ra, ta). This step increases handwriting style diversity for those classes without requiring new data collection.

D. Step 3: Structure-Preserving Augmentation

This is the primary methodological contribution of the study. Augmentation is applied *exclusively to the 100 non-A classes* until each class reaches a target of $T = 300$ training images. New images are created by repeatedly sampling from existing training images and applying transformations until each class reaches T samples.

TABLE II
STRUCTURE-PRESERVING AUGMENTATION STRATEGY

Category	Transformation	Status
Photometric	Brightness jitter ($\pm 30\%$)	✓
	Contrast jitter ($\pm 30\%$)	✓
	Gaussian blur (radius 0.5–1.5 px)	✓
	Gaussian noise ($\sigma = 1\text{--}5\%$)	✓
	Sharpening ($1.2\times\text{--}2.0\times$)	✓
Morphological	Dilation (MaxFilter $3\times 3 / 5\times 5$)	✓
	Erosion (MinFilter $3\times 3 / 5\times 5$)	✓
Spatial	Random crop + resize (max 10% margin)	✓
Geometric (prohibited)	Rotation ($> \pm 3^\circ$)	×
	Horizontal / vertical flip	×
	Shear / affine transform	×
	Elastic distortion	×

The most critical design constraint is the prohibition of geometric transformations in any form. Wisesty et al. [3] empirically demonstrated that augmentation actually degrades Javanese script recognition because transformations alter visual characteristics that distinguish between classes. This is reinforced by the finding of Nindya et al. [6] that the “ha” and “la” pair are nearly identical horizontal reflections, so applying *horizontal flip* would directly create cross-class ambiguous labels.

Table II lists the permitted and prohibited transformations. Permitted transformations target *appearance variability* (ink density, scan noise, pen stroke thickness) without altering the spatial layout of the character. Each augmented image is created by applying 1–3 transformations selected randomly from the permitted list. Augmentation is performed *offline* (saved to disk) to ensure reproducibility. The final training set after augmentation totals 38,510 images, an increase of approximately $4.0\times$ from 9,549 images before augmentation (8,483 original training images plus 1,066 variation images).

E. Step 4: Weighted Loss Function

After augmentation, class weights are computed using:

$$w_c = \frac{N}{C \cdot n_c} \quad (1)$$

where N is the total training samples, $C = 120$ is the number of classes, and n_c is the number of training samples for class c . This formula is equivalent to `compute_class_weight('balanced')` in scikit-learn, which assigns greater loss penalties to misclassification of minority classes. These weights are applied as a parameter to the `CrossEntropyLoss` function during training.

F. Model Architectures

Once the data pipeline is established, the next stage is selecting the model architectures to evaluate. Four CNN ar-

TABLE III
CNN ARCHITECTURES COMPARED

Architecture	Parameters	Literature Relevance
EfficientNet-B0 [13]	~5.3M	—
MobileNetV3-Large [10]	~5.4M	Wisesty et al. [3]
ResNet-18 [12]	~11.7M	Ifriza et al. [5]
VGG-16 [11]	~138M	Nindya et al. [6]

TABLE IV
HYPERPARAMETER CONFIGURATION PER ARCHITECTURE

Architecture	LR Phase 1	LR Phase 2	WD	ES
EfficientNet-B0	10^{-3}	10^{-4}	10^{-5}	10
MobileNetV3-Large	10^{-3}	10^{-4}	10^{-6}	8
ResNet-18	2×10^{-3}	2×10^{-4}	10^{-5}	12
VGG-16	5×10^{-4}	2×10^{-5}	10^{-4}	15

WD = weight decay. ES = early stopping patience.

chitectures with ImageNet-1K pretrained weights are selected based on two complementary criteria. The first criterion is representation of conceptually distinct architecture families: compound scaling efficiency (EfficientNet), residual connections (ResNet), *depthwise separable convolution* (MobileNetV3), and deep convolutional blocks (VGG). The second criterion is direct relevance to prior Javanese script research, enabling fair comparison against architectures already reported in the literature. Table III summarizes the four architectures and their literature relevance. Each architecture’s classification head is replaced with a linear layer mapping to 120 classes, while all backbone weights are initialized from ImageNet-1K pretraining.

G. Training Strategy

All models follow a uniform two-phase fine-tuning strategy to ensure fair cross-architecture comparison. In Phase 1 over 10 epochs, all backbone parameters are frozen and only the classification head is trained using the Adam optimizer with a higher learning rate. This phase aims to build a stable head initialization before gradients propagate to the more sensitive pretrained feature layers. In Phase 2 running up to 90 epochs, all layers are unfrozen for full fine-tuning with a lower learning rate. During Phase 2, a ReduceLROnPlateau scheduler with factor 0.5 and *patience* 3 adaptively reduces the learning rate when validation loss stagnates, while *early stopping* halts training if no improvement occurs for a number of epochs as specified in Table IV. All models are trained using the weighted CrossEntropyLoss per (1). Full hyperparameter configuration per architecture is shown in Table IV.

All experiments were run on an NVIDIA GeForce RTX 3070 Ti GPU (8GB VRAM) with an AMD Ryzen 5 3600 CPU, using Python 3.10, PyTorch 2.9.1 (CUDA 13.0), torchvision 0.24.1, and scikit-learn. All random seeds were fixed at 42

for reproducibility, with a batch size of 32 and 3 DataLoader workers.

Evaluation is performed on the test set (1,919 samples) using Top-1 accuracy, macro F1-score, and weighted F1-score as overall metrics. Macro F1-score is chosen as the primary metric because it assigns equal weight to each class regardless of majority class dominance. As the primary analytical contribution, macro F1-score is also computed separately for the Vowel A group (20 majority classes, approximately 1,519 test samples) and the non-A vowel group (100 minority classes, 400 test samples), with the absolute difference between the two used to measure model generalization balance.

V. RESULTS AND DISCUSSION

A. Overall Performance Comparison

Table V summarizes the performance of all architectures on the test set. EfficientNet-B0 achieves the highest accuracy of 97.08% with a macro F1-score of 0.9110, while ResNet-18 achieves the highest macro F1-score (0.9180). MobileNetV3-Large and VGG-16 show lower performance despite having 5.4 and 138 million parameters respectively, indicating that parameter count alone is not deterministic of performance under these imbalanced conditions.

B. Vowel A vs. Non-A Vowel Gap Analysis

The Gap column in Table V reveals two important patterns hidden by aggregate metrics, which constitute the primary analytical contribution of this study.

EfficientNet-B0 exhibits a pattern typical of imbalanced conditions, where the very high Vowel A F1 (0.9902) contrasts with a lower non-A Vowel F1 (0.8783), yielding a gap of 0.1119. This confirms that despite high aggregate accuracy, the model retains residual bias toward majority classes not fully eliminated by the augmentation pipeline and weighted loss function.

The most surprising finding comes from ResNet-18, which shows an F1 gap between Vowel A and non-A Vowel classes of only 0.0004 (0.8972 vs. 0.8976) — a difference that is practically negligible. This was unexpected because ResNet-18 is not the highest-accuracy architecture, yet it is the architecture that most successfully generalizes in a balanced manner to minority classes. This capability likely stems from the *residual connection* mechanism, which allows gradients to flow more stably to early layers, so low-level visual features relevant to minority script characters — which have only approximately 300 augmented training images — can still be learned without distortion from majority-class gradient dominance. Such findings would be invisible in prior studies using balanced datasets, because the majority-minority gap was never measured as a variable.

VGG-16 and MobileNetV3-Large exhibit larger gaps of 0.2422 and 0.3393. For VGG-16, the large gap despite relatively high overall accuracy (93.38%) reflects an architecture more prone to over-optimizing majority classes. MobileNetV3-Large, which was the best-performing choice of Wisesty et

TABLE V
PERFORMANCE COMPARISON AND GAP ANALYSIS OF FOUR ARCHITECTURES ON THE 120-CLASS TEST SET

Architecture	Accuracy (%)	Macro F1	Vowel A F1	Non-A F1	Gap	Epochs
EfficientNet-B0	97.08	0.9110	0.9902	0.8783	0.1119	50
ResNet-18	96.93	0.9180	0.8972	0.8976	0.0004	36
VGG-16	93.38	0.7867	0.9824	0.7402	0.2422	36
MobileNetV3-Large	90.83	0.6935	0.9714	0.6321	0.3393	52

Gap = |Vowel A F1 − Non-A F1|.




Fig. 2. EfficientNet-B0 training curves. The Phase 1-to-Phase 2 transition is visible at epoch 11.




Fig. 3. ResNet-18 training curves. Faster convergence compared to EfficientNet-B0 (36 vs. 50 epochs).

al. [3] on a balanced 20-class dataset, shows the weakest performance among the four architectures under these imbalanced 120-class conditions.

C. Training Curve Analysis

Fig. 2 and 3 display the training curves of the two best architectures. The Phase 1 to Phase 2 transition (epoch 11) is clearly visible as a sharp loss drop, confirming the effectiveness of the two-phase fine-tuning strategy. EfficientNet-B0 converges at epoch 50 (validation accuracy 97.60%), while ResNet-18 converges earlier at epoch 36, reflecting better training efficiency.

D. Comparison with Prior Work

Table VI contextualizes this study within the existing Javanese script recognition literature. This study is the first to evaluate all 120 syllabic classes on a naturally imbalanced

TABLE VI
COMPARISON WITH JAVANESE SCRIPT RECOGNITION STUDIES

Study	Classes	Balanced?	Best Metric
Susanto et al. [9]	120	Yes	Accuracy 97.29%
Susanto et al. [8]	120	Yes	Accuracy 99.65%
Ifriza et al. [5]	20	Yes	Accuracy 95.53%
Wisesty et al. [3]	20	Yes	F1 = 0.979
Nindya et al. [6]	20	Yes	Accuracy 93.19%
Faizin et al. [2]	20	No	–
This work	120	No	Accuracy 97.08%

dataset, and the first to report four-architecture comparisons with per-vowel-group analysis.

EfficientNet-B0 achieves accuracy comparable to prior balanced studies (97.08% vs. 97.29% by Susanto et al. [9]), despite being trained on a dataset with 33.10× imbalance. This confirms that the proposed pipeline successfully bridges naturally imbalanced conditions to near the performance of ideal balanced conditions. Furthermore, no prior study has reported per-vowel-group analysis, making the finding of ResNet-18’s generalization gap of only 0.0004 an entirely novel contribution to this literature.

One limitation must be acknowledged: each non-A Vowel class is represented by only 4 images in the test set, so a single misprediction already shifts per-class F1 by 25%. Therefore, all conclusions are based on per-vowel-group macro F1 (averaged over 20 classes), and larger minority data collection remains a priority for follow-up research.

VI. CONCLUSION

This study demonstrates that 120-class handwritten Javanese script recognition under severe natural class imbalance can be achieved at competitive performance levels through a combination of a structure-preserving data preparation pipeline and transfer learning. EfficientNet-B0 achieves 97.08% accuracy comparable to prior studies using artificially balanced datasets [9], proving that imbalanced conditions are not an insurmountable obstacle given an appropriate data strategy.

The analytically most significant finding is the performance of ResNet-18, which exhibits an F1 gap between majority and

minority classes of only 0.0004 — a generalization balance that has never been reported in the Javanese script literature because no prior study has measured this dimension. This finding has direct practical implications: when a Javanese OCR system is deployed on real documents spanning the full vowel spectrum uniformly, ResNet-18 is the more appropriate choice than architectures with higher aggregate accuracy but bias toward majority classes.

This study is limited to isolated character recognition and a small test set for minority classes. Going forward, collecting more minority data will strengthen the validation of these findings, while exploring GAN-based augmentation [2] can complement the proposed photometric approach. Extension to word-level recognition incorporating *pasangan* forms represents the most application-relevant next step [1].

REFERENCES

- [1] B. A. Wicaksono, B. S. Hantono, and T. B. Adji, “Word segmentation task for Southeast Asian abugida scripts: A systematic literature review,” in *Proc. 2024 2nd Int. Conf. Technology Innovation and Its Applications (ICTIIA)*, 2024.
- [2] M. A. Faizin, N. Suciati, and C. Fatichah, “Handling imbalance in Javanese manuscript character dataset using skeleton-based balancing generative adversarial networks,” *Jurnal RESTI (Rekayasa Sistem dan Teknologi Informasi)*, vol. 9, no. 4, pp. 820–831, 2025.
- [3] U. N. Wisesty et al., “MNetNCR: MobileNet model for efficient traditional Nusantara script character recognition,” *International Journal of Artificial Intelligence (IJ-AI)*, vol. 15, no. 2, pp. 1513–1528, 2026.
- [4] R. Widiarti and F. T. Adji, “A baseline evaluation of OCR segmentation and classification methods for printed Javanese script,” *Engineering, Technology & Applied Science Research (ETASR)*, vol. 16, no. 1, pp. 31699–31705, 2026.
- [5] Y. N. Ifriza et al., “Optimizing Javanese script recognition using fine-tuned ResNet-18 and transfer learning,” *International Journal of Artificial Intelligence (IJ-AI)*, vol. 15, no. 1, pp. 443–453, 2026.
- [6] T. P. Nindya, M. D. Sulistiyono, and U. N. Wisesty, “Recognition of Balinese, Javanese, and Sundanese scripts using CNN with VGG-16 architecture,” in *Proc. 2025 Int. Conf. Data Science and Its Applications (ICoDSA)*, 2025, pp. 1221–1226.
- [7] M. D. Sulistiyono et al., “Aksarawa: Comparative study of shallow learning for OCR of Nusantara scripts,” *Jurnal RESTI*, vol. 9, no. 6, pp. 1537–1546, 2025.
- [8] A. Susanto et al., “Handwritten Javanese script recognition method based 12-layers deep convolutional neural network and data augmentation,” *International Journal of Artificial Intelligence (IJ-AI)*, vol. 12, no. 3, pp. 1448–1458, 2023.
- [9] A. Susanto, I. U. W. Mulyono, C. A. Sari, E. H. Rachmawanto, and D. R. I. M. Setiadi, “Improved Javanese script recognition using custom model of convolution neural network,” *International Journal of Electrical and Computer Engineering (IJECE)*, vol. 13, no. 6, pp. 6629–6636, 2023.
- [10] A. Howard et al., “Searching for MobileNetV3,” in *Proc. IEEE/CVF Int. Conf. Computer Vision (ICCV)*, 2019, pp. 1314–1324.
- [11] K. Simonyan and A. Zisserman, “Very deep convolutional networks for large-scale image recognition,” in *Proc. Int. Conf. Learning Representations (ICLR)*, 2015.
- [12] K. He, X. Zhang, S. Ren, and J. Sun, “Deep residual learning for image recognition,” in *Proc. IEEE Conf. Computer Vision and Pattern Recognition (CVPR)*, 2016, pp. 770–778.
- [13] M. Tan and Q. V. Le, “EfficientNet: Rethinking model scaling for convolutional neural networks,” in *Proc. Int. Conf. Machine Learning (ICML)*, 2019, pp. 6105–6114.